

FINAL PROJECT REPORT – Pneumonia detection challenge

Batch name- AIML online November 19-B

GROUP-4 AIML CAPSTONE -CV PROJECT

Done by

Gautam Ravi kumar

Shruthi V

Aishvarya Shri

Piyush Panchariya

Akshay Badhani

RSNA PNEUMONIA DETECTION CHALLENGE

Abstract

Pneumonia is a condition caused due to the inflammation of the lungs that primarily affects the small type of air sacs called as alveoli. This disease can be caused by bacteria, fungi or virus. It is often diagnosed by symptoms, physical type of examination and by employing the method of chest radiograph (CXR). Factors such as positioning of the patient and the depth of inspiration result in changing the appearance of the CXR which leads to complicated interpretations further. The objective of this study is to build an algorithm to detect a visual signal for pneumonia in the medical images that specifically locate the lung opacities and the position of inflammation in an image on the chest radiographs. The data set consists of the medical images which contain a combination of the header meta data as well as raw image arrays for the pixel data. The progress of the study includes pre-processing, data visualization and exploratory data analysis of the given data set. Then an object detection model is built and employment of transfer learning process to fine tune the model. Thus, based on the analysis made, certain automation systems can be made possible in detection of pneumonia which will help doctors to make better clinical decisions using the powerful AI techniques.

Pneumonia Detection

Pneumonia accounts for over 15% of all deaths of children under 5 years old internationally. In 2017, 920,000 children under the age of 5 died from the disease. It requires review of a chest radiograph (CXR) by highly trained specialists and confirmation through clinical history, vital signs and laboratory exams. Pneumonia usually manifests as an area or areas of increased opacity on CXR. However, the diagnosis of pneumonia on CXR is complicated because of a number of other conditions in the lungs such as fluid overload (pulmonary edema), bleeding, volume loss (atelectasis or collapse), lung cancer, or post-radiation or surgical changes. Outside of the lungs, fluid in the pleural space (pleural effusion) also appears as increased opacity on CXR. When available, comparison of CXRs of the patient taken at different time points and correlation with clinical symptoms and history are helpful in making the diagnosis.

Problem statement

Pneumonia accounts for the increase in the mortality rate of the human beings which demands for this AI challenge towards detection and localising the pneumonia. Factors pertaining to the complications in the interpretations made by the doctors by looking at the images of the CXR are change in the positioning of the patient and the depth of the inspiration can alter the appearance of the CXR images. Because of this, doctors would face difficulties in treating the patients since they will be faced with high volumes of images coming day by day during every shifts. The exact position of the inflammation of the lungs has to be shown in the image by detecting the presence of lung opacities. It is also given that all the lung opacities will not be a sign of pneumonia as given in the data set to be “not normal no lung opacity”. This denotes some kind of other abnormalities in the lungs apart from pneumonia. Hence, a robust AI technology for the detection of inflammation in the lungs is necessary.

Objectives of the study

1. To perform Data pre-processing, EDA and Visualisation of the data
2. To construct a pneumonia detection model (Object detection system)
3. To Test the model and to fine tune the processes to improve accuracy using transfer learning.

Data Collection

Data sets are downloaded from the Kaggle website in a competition organised by the Radiological Society of North America (RSNA) pneumonia detection challenge which consists of the train and test data of CRX images. The data set consists of the Patient id, lung x ray images and class of the degree of presence of lung opacities.

Exploratory Data Analysis (EDA)

1. Read the data into a notebook
2. Import of the necessary libraries
3. To find the shape of the data and type of individual columns
4. Descriptive statistics of the attributes in the data
5. Check the presence of missing values and unique values
6. Data Visualization and distribution of the categorical columns

Several key items in the data set given:

1. Stage_2_detailed_class_info.csv: CSV file containing the training set patient ids and labels of the class.
2. stage_2_sample_submission.csv: CSV file containing the patient ids and the prediction strings.
3. stage_2_train_labels.csv: CSV file containing the patient ids with x,y, width, height and targets involved.
4. Stage_2_test images and Stage_2_train images: Contains Medical X ray images of the lungs which is a DCM file type

Data Description

```
#Reading the data using pandas
trainlabels=pd.read_csv(r"C:\Users\User\Desktop\stage_2_train_labels.csv")
trainlabels.head()
print(trainlabels.iloc[0])
```

```
patientId    0004cfab-14fd-4e49-80ba-63a80b6bddd6
x                                                    NaN
y                                                    NaN
width                                                NaN
height                                               NaN
Target                                              0
Name: 0, dtype: object
```

From the above, we can see that each row in the CSV file contains a patient Id (one unique value per patient), a target (either 0 or 1 for absence or presence of pneumonia, respectively) and the corresponding abnormality bounding box defined by the upper-left hand corner (x, y) coordinate and its corresponding width and height. In this particular case, the patient does *not* have pneumonia and so the corresponding bounding box information is set to NaN.

The below code represents a patient with pneumonia as the target of 1 implies

```
print(trainlabels.iloc[4])

patientId    00436515-870c-4b36-a041-de91049b9ab4
x                                                    264
y                                                    152
width                                                213
height                                               379
Target                                              1
Name: 4, dtype: object
```

Null values and information about the data

```
In [222]: class_info.head(10)
```

```
Out[222]:
```

| | patientId | class |
|---|---------------------------------------|------------------------------|
| 0 | 0004cfab-14fd-4e49-80ba-63a80b6bdddd6 | No Lung Opacity / Not Normal |
| 1 | 00313ee0-9eaa-42f4-b0ab-c148ed3241cd | No Lung Opacity / Not Normal |
| 2 | 00322d4d-1c29-4943-afc9-b6754be640eb | No Lung Opacity / Not Normal |
| 3 | 003d8fa0-6bf1-40ed-b54c-ac657f8495c5 | Normal |
| 4 | 00436515-870c-4b36-a041-de91049b9ab4 | Lung Opacity |
| 5 | 00436515-870c-4b36-a041-de91049b9ab4 | Lung Opacity |
| 6 | 00569f44-917d-4c86-a842-81832af98c30 | No Lung Opacity / Not Normal |
| 7 | 006cec2e-6ce2-4549-bffa-eadfd1e9970 | No Lung Opacity / Not Normal |
| 8 | 00704310-78a8-4b38-8475-49f4573b2dbb | Lung Opacity |
| 9 | 00704310-78a8-4b38-8475-49f4573b2dbb | Lung Opacity |

```
In [223]: class_info.isnull().sum()
```

```
Out[223]: patientId    0
          class        0
          dtype: int64
```

There are no null values present in the data and the class corresponding to the patient id shows whether they have pneumonia or not.

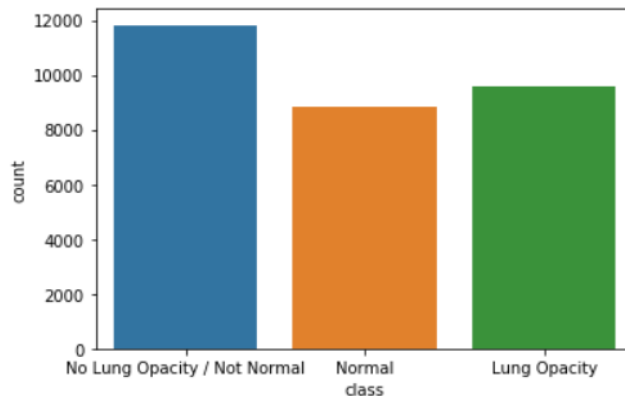
Normal- No defect in the lungs (Absence of pneumonia)

Lung opacity – Lungs affected by pneumonia

No lung opacity/Not Normal – Defective lung but not by pneumonia (may be due to some other abnormalities in the lungs)

Count plot for the number of cases:

```
In [224]: sns.countplot(x='class', data=class_info)
Out[224]: <matplotlib.axes._subplots.AxesSubplot at 0x13b99bd5f48>
```



```
In [225]: class_info['class'].value_counts(dropna=False)
Out[225]: No Lung Opacity / Not Normal    11821
Lung Opacity                            9555
Normal                                  8851
Name: class, dtype: int64
```

From the plot, it is evident that there are 9555 people affected by pneumonia and 11821 people affected by some other abnormalities in the lungs.

Merge of the data sets

The two csv files `stage_2_trainlabels` and `stage_2_detailedclass_info` are merged together based on the patient id as a criteria as given below. In detailed class info dataset are given the detailed information about the type of positive or negative class associated with a certain patient. In Train labels dataset are given the patient ID and the window (x min, y min, width and height of the) containing evidence of pneumonia.

```
In [235]: # Let's merge now the two datasets, using Patient ID as the merge criteria.
train_merged = train_labels.merge(class_info, left_on='patientId',right_on='patientId', how='inner')
train_merged=train_merged.drop_duplicates()
train_merged.reset_index(drop=True, inplace=True)
train_merged
```

```
Out[235]:
```

| | patientId | x | y | width | height | Target | class |
|-------|--------------------------------------|-------|-------|-------|--------|--------|------------------------------|
| 0 | 0004cfab-14fd-4e49-80ba-63a80b6bdd6 | NaN | NaN | NaN | NaN | 0 | No Lung Opacity / Not Normal |
| 1 | 00313ee0-9eaa-42f4-b0ab-c148ed3241cd | NaN | NaN | NaN | NaN | 0 | No Lung Opacity / Not Normal |
| 2 | 00322d4d-1c29-4943-afc9-b6754be640eb | NaN | NaN | NaN | NaN | 0 | No Lung Opacity / Not Normal |
| 3 | 003d8fa0-6bf1-40ed-b54c-ac657f8495c5 | NaN | NaN | NaN | NaN | 0 | Normal |
| 4 | 00436515-870c-4b36-a041-de91049b9ab4 | 264.0 | 152.0 | 213.0 | 379.0 | 1 | Lung Opacity |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 30222 | c1ec14ff-f6d7-4b38-b0cb-fe07041cbdc8 | 185.0 | 298.0 | 228.0 | 379.0 | 1 | Lung Opacity |
| 30223 | c1edf42b-5958-47ff-a1e7-4f23d99583ba | NaN | NaN | NaN | NaN | 0 | Normal |
| 30224 | c1f6b555-2eb1-4231-98f6-50a963976431 | NaN | NaN | NaN | NaN | 0 | Normal |
| 30225 | c1f7889a-9ea9-4acb-b64c-b737c929599a | 570.0 | 393.0 | 261.0 | 345.0 | 1 | Lung Opacity |
| 30226 | c1f7889a-9ea9-4acb-b64c-b737c929599a | 233.0 | 424.0 | 201.0 | 356.0 | 1 | Lung Opacity |

30227 rows × 7 columns

The merged data set is grouped together based on the targets given to the patients as 0 and 1 where 0 is for No lung opacity/ Not normal (Other abnormalities) and 1 is for lung opacity which shows the sign for existence of pneumonia. The columns are grouped as follows.

```
In [227]: temporary = train_merged.groupby('Target')['class'].value_counts()
temporary
```

```
Out[227]: Target  class
0      No Lung Opacity / Not Normal    11821
      Normal                        8851
1      Lung Opacity                    9555
Name: class, dtype: int64
```

All chest examinations with Target = 1 (pathology detected) associated with class: Lung Opacity. The chest examinations with Target = 0 (no pathology detected) are either of class: Normal or class: No Lung Opacity / Not Normal.

Overview of the DICOM files and Medical Images

Medical images are stored in a special format known as DICOM files (*.dcm). These type of files contain a combination of header meta data along with the underlying raw image arrays for the pixel data available. They contain a combination of header metadata as well as underlying raw image arrays for pixel data. In Python, one popular library to access and manipulate DICOM files is the pydicom module. To use the pydicom library, we have to first find out the DICOM file for a given patient Id by simply looking for the matching file in the stage_2_train_images file, and the use the pydicom.read_file() method to load the data.

```
dcm_data = pydicom.dcmread(train_img[0])
dcm_data
```

```

Out[254]: Dataset.file_meta -----
(0002, 0000) File Meta Information Group Length  UL: 202
(0002, 0001) File Meta Information Version       OB: b'\x00\x01'
(0002, 0002) Media Storage SOP Class UID        UI: Secondary Capture Image Storage
(0002, 0003) Media Storage SOP Instance UID     UI: 1.2.276.0.7230010.3.1.4.8323329.28530.1517874485.775526
(0002, 0010) Transfer Syntax UID               UI: JPEG Baseline (Process 1)
(0002, 0012) Implementation Class UID          UI: 1.2.276.0.7230010.3.0.3.6.0
(0002, 0013) Implementation Version Name       SH: 'OFFIS_DCMTK_360'
-----
(0008, 0005) Specific Character Set             CS: 'ISO_IR 100'
(0008, 0016) SOP Class UID                     UI: Secondary Capture Image Storage
(0008, 0018) SOP Instance UID                  UI: 1.2.276.0.7230010.3.1.4.8323329.28530.1517874485.775526
(0008, 0020) Study Date                       DA: '19010101'
(0008, 0030) Study Time                       TM: '000000.00'
(0008, 0050) Accession Number                 SH: ''
(0008, 0060) Modality                         CS: 'CR'
(0008, 0064) Conversion Type                  CS: 'WSD'
(0008, 0090) Referring Physician's Name       PN: ''
(0008, 103e) Series Description                LO: 'view: PA'
(0010, 0010) Patient's Name                   PN: '0004cfab-14fd-4e49-80ba-63a80b6bddd6'
(0010, 0020) Patient ID                       LO: '0004cfab-14fd-4e49-80ba-63a80b6bddd6'
(0010, 0030) Patient's Birth Date             DA: ''
(0010, 0040) Patient's Sex                    CS: 'F'
(0010, 1010) Patient's Age                    AS: '51'
(0018, 0015) Body Part Examined               CS: 'CHEST'
(0018, 5101) View Position                    CS: 'PA'
(0020, 000d) Study Instance UID                UI: 1.2.276.0.7230010.3.1.2.8323329.28530.1517874485.775525
(0020, 000e) Series Instance UID              UI: 1.2.276.0.7230010.3.1.3.8323329.28530.1517874485.775524
(0020, 0010) Study ID                         SH: ''
(0020, 0011) Series Number                    IS: "1"
(0020, 0013) Instance Number                  IS: "1"
(0020, 0020) Patient Orientation              CS: ''
(0028, 0002) Samples per Pixel                US: 1
(0028, 0004) Photometric Interpretation       CS: 'MONOCHROME2'
(0028, 0010) Rows                            US: 1024
(0028, 0011) Columns                          US: 1024
(0028, 0030) Pixel Spacing                    DS: [0.14300000000000002, 0.14300000000000002]
(0028, 0100) Bits Allocated                   US: 8

(0010, 0020) Patient ID                       LO: '0004cfab-14fd-4e49-80ba-63a80b6bddd6'
(0010, 0030) Patient's Birth Date             DA: ''
(0010, 0040) Patient's Sex                    CS: 'F'
(0010, 1010) Patient's Age                    AS: '51'
(0018, 0015) Body Part Examined               CS: 'CHEST'
(0018, 5101) View Position                    CS: 'PA'
(0020, 000d) Study Instance UID                UI: 1.2.276.0.7230010.3.1.2.8323329.28530.1517874485.775525
(0020, 000e) Series Instance UID              UI: 1.2.276.0.7230010.3.1.3.8323329.28530.1517874485.775524
(0020, 0010) Study ID                         SH: ''
(0020, 0011) Series Number                    IS: "1"
(0020, 0013) Instance Number                  IS: "1"
(0020, 0020) Patient Orientation              CS: ''
(0028, 0002) Samples per Pixel                US: 1
(0028, 0004) Photometric Interpretation       CS: 'MONOCHROME2'
(0028, 0010) Rows                            US: 1024
(0028, 0011) Columns                          US: 1024
(0028, 0030) Pixel Spacing                    DS: [0.14300000000000002, 0.14300000000000002]
(0028, 0100) Bits Allocated                   US: 8
(0028, 0101) Bits Stored                      US: 8
(0028, 0102) High Bit                         US: 7
(0028, 0103) Pixel Representation             US: 0
(0028, 2110) Lossy Image Compression          CS: '01'
(0028, 2114) Lossy Image Compression Method  CS: 'ISO_10918_1'
(7fe0, 0010) Pixel Data                       OB: Array of 142006 elements

```

We can observe that we do have available some useful information in the DICOM metadata with predictive value such as Patient sex, Patient age, Modality, Body part examined, Positional view, Rows and columns and Pixel spacing.

Most of the standard headers containing patient identifiable information have been anonymized (removed) so we are left with a relatively sparse set of metadata.

Exploration of the DICOM data

Let's check how many images are in the train and test folders containing the images of the lungs of the patients.

```
In [229]: train_img=sorted(glob.glob(path+'/stage_2_train_images/*'))
test_img=sorted(glob.glob(path+'/stage_2_test_images/*'))
print('No. of train images:',len(train_img))
print('First train image filename:',train_img[0])
print('No. of test images:',len(test_img))
print('First 5 test image name:',test_img[0])
```

No. of train images: 26684
First train image filename: C:/Users/Dell/Desktop/Capstone/Dataset/stage_2_train_images\0004cfab-14fd-4e49-80ba-63a80b6bdd6.dcm
No. of test images: 3000
First 5 test image name: C:/Users/Dell/Desktop/Capstone/Dataset/stage_2_test_images\0000a175-0e68-4ca4-b1af-167204a7e0bc.dcm

There are 26684 images in the train image set and 3000 images present in the test image set.

To check the presence of duplicates in the train data set

```
In [228]: #number of unique patients outta overall sample
train_merged['patientId'].nunique()
```

Out[228]: 26684

We confirmed that the number of *unique* patients Id are equal with the number of DICOM images in the train set. It is evident that there are 26684 patient ids which are equal to the number of train images as given as explained from the code above.

DCM image is plotted now using the age, sex and ID of the patient as given below.

```
In [193]: ax=plt.gca()
pid=dc_m_data.PatientID
age = dc_m_data.PatientAge
sex = dc_m_data.PatientSex
plt.imshow(im)
plt.axis("off")
ax.set_title('ID: {}\nAge: {} Sex: {}'.format(pid,age, sex))
```

Out[193]: Text(0.5, 1.0, 'ID: 00704310-78a8-4b38-8475-49f4573b2dbb\nAge: 75 Sex: M')

ID: 00704310-78a8-4b38-8475-49f4573b2dbb
Age: 75 Sex: M



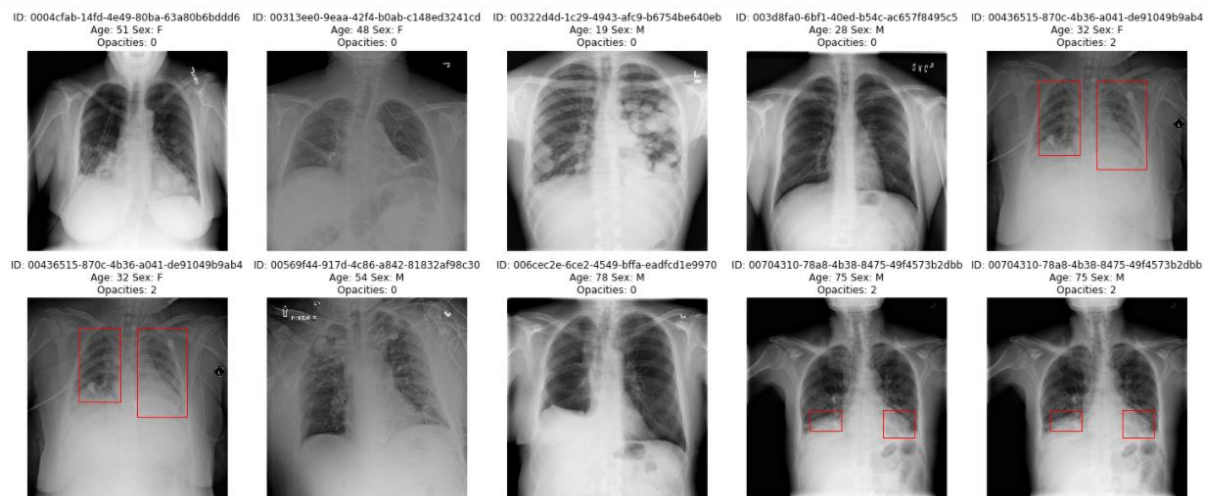
Then the merged data set consisting of the classes of abnormalities in the lungs is given as the criteria in the differentiation of the dicom images of the patients as given in the code below which denotes the number of lung opacities in the patients through the images. These are based on the Patient id, age, sex and the classes.

```
In [195]: fig=plt.figure(figsize=(25, 12))
for i in range(0,10):
    dc_m_data = pydicom.dcmread(train_img[i])
    im=dc_m_data.pixel_array
    fig.add_subplot(2, 5, i+1)
    ax=plt.gca()
    stat=train_merged[train_merged.patientId==dc_m_data.PatientID]["Target"].sum()
    pid=dc_m_data.PatientID
    age = dc_m_data.PatientAge
    sex = dc_m_data.PatientSex
    plt.imshow(im)
    plt.axis("off")
    ax.set_title('ID: {}\nAge: {} Sex: {}\nOpacities: {}'.format(pid,age,sex,stat))
```



Plot of the rectangular bounding boxes for the existence of the lung opacities which shows a target of 1.

```
In [161]: fig=plt.figure(figsize=(25, 10))
for i in range(10):
    patientId = df_t0['patientId'][i]
    train_img=path+'/stage_2_train_images/%s.dcm'%patientId
    dcm_data = pydicom.dcmread(train_img)
    im=dcm_data.pixel_array
    fig.add_subplot(2, 5, i+1)
    plt.imshow(im)
    plt.axis("off")
    ax = plt.gca()
    df_temp=train_labels[train_labels['patientId'] == patientId]
    for index, row in df_temp.iterrows():
        rect = plt.Rectangle((row['x'], row['y']), row['width'], row['height'], color='red',fill=False)
    ax.add_patch(rect)
```



Building up a model

Mask-RCNN model is employed for the detection of pneumonia along with the process of transfer learning. The type of weights that are used are COCO weights which is a large scale object detection, segmentation and captioning of the data sets. The pre trained COCO weights are used in this process.

Mask- RCNN model

Mask RCNN is a type of deep neural network which is designed and aimed at solving segmentation problems in computer vision as well as machine learning. It is used for the separation of different types of objects in any given type of an image or a video. In case of an image, this model helps in giving out the object bounding boxes, classes and the masks. In this case this model is used for the identification of the extent of inflammation in the lungs and the other sorts of abnormalities with the help of the bounding boxes as a part of identification and segmentation modelling.

There are two stages present in the Mask RCNN model. First step is that it helps in the generation of the proposals about the presence of regions where there may be an object shown based on the input image given. The next step would be, it will also predict the class of the object and refines the bounding boxes which leads to the generation of the mask in the level of pixels of the detected object based on the first step and both the steps would be interconnected to each other.

The weights for the pre trained model is downloaded in which the Mask R-CNN model is trained on the MS coco Datasets. Finally, we will use the Mask R-CNN architecture and the pretrained weights to generate predictions for our own images.

```
In [ ]: # Import Mask RCNN
sys.path.append(os.path.join(work_dir, 'Mask_RCNN')) # To find local version of the library
from mrcnn.config import Config
from mrcnn import utils
import mrcnn.model as modellib
from mrcnn import visualize
from mrcnn.model import log

In [ ]: # Downloading COCO pretrained weight
!wget --quiet https://github.com/matterport/Mask_RCNN/releases/download/v2.0/mask_rcnn_coco.h5
!ls -lh mask_rcnn_coco.h5

weights = "mask_rcnn_coco.h5"
```

On experimenting with 2 epochs, we got the following results,

Rpn_class_loss : 0.028

Rpn_bbox_loss : 0.5910

Mrcnn_class_loss: 0.2778

Mrcnn_bbox_loss: 0.5423

Mrcnn_mask_loss: 0.3967

We could observe that the bounding box loss with masked RCNN is around 0.5 and has to be improved further. This loss depicts the difference in the bounding boxes that are predicted to the ground truth.

Resnet Model

Residual Network (ResNet) is a Convolutional Neural Network (CNN) architecture which was designed to enable hundreds or thousands of convolutional layers. While previous CNN architectures had a drop off in the effectiveness of additional layers, ResNet can add a large number of layers with strong performance.

ResNet was an innovative solution to the “vanishing gradient” problem. Neural networks train via the backpropagation process (see our guide on backpropagation), which relies on gradient descent, moving down the loss function to find the weights that minimize it. If there are too many layers, repeated multiplication makes the gradient smaller and smaller, until it “disappears”, causing performance to saturate or even degrade with each additional layer.

The ResNet solution is “identity shortcut connections”. ResNet stacks up identity mappings, layers that initially don’t do anything, and skips over them, reusing the activations from previous layers. Skipping initially compresses the network into only a few layers, which enables faster learning. Then, when the network trains again, all layers are expanded and the “residual” parts of the network explore more and more of the feature space of the source image.



model.summary()



Model: "functional_1"

| Layer (type) | Output Shape | Param # | Connected to |
|--|-----------------------|---------|--|
| input_1 (InputLayer) | [(None, 224, 224, 3)] | 0 | |
| conv1_pad (ZeroPadding2D) | (None, 230, 230, 3) | 0 | input_1[0][0] |
| conv1_conv (Conv2D) | (None, 112, 112, 64) | 9472 | conv1_pad[0][0] |
| conv1_bn (BatchNormalization) | (None, 112, 112, 64) | 256 | conv1_conv[0][0] |
| conv1_relu (Activation) | (None, 112, 112, 64) | 0 | conv1_bn[0][0] |
| pool1_pad (ZeroPadding2D) | (None, 114, 114, 64) | 0 | conv1_relu[0][0] |
| pool1_pool (MaxPooling2D) | (None, 56, 56, 64) | 0 | pool1_pad[0][0] |
| conv2_block1_1_conv (Conv2D) | (None, 56, 56, 64) | 4160 | pool1_pool[0][0] |
| conv2_block1_1_bn (BatchNormalization) | (None, 56, 56, 64) | 256 | conv2_block1_1_conv[0][0] |
| conv2_block1_1_relu (Activation) | (None, 56, 56, 64) | 0 | conv2_block1_1_bn[0][0] |
| conv2_block1_2_conv (Conv2D) | (None, 56, 56, 64) | 36928 | conv2_block1_1_relu[0][0] |
| conv2_block1_2_bn (BatchNormalization) | (None, 56, 56, 64) | 256 | conv2_block1_2_conv[0][0] |
| conv2_block1_2_relu (Activation) | (None, 56, 56, 64) | 0 | conv2_block1_2_bn[0][0] |
| conv2_block1_0_conv (Conv2D) | (None, 56, 56, 256) | 16640 | pool1_pool[0][0] |
| conv2_block1_3_conv (Conv2D) | (None, 56, 56, 256) | 16640 | conv2_block1_2_relu[0][0] |
| conv2_block1_0_bn (BatchNormalization) | (None, 56, 56, 256) | 1024 | conv2_block1_0_conv[0][0] |
| conv2_block1_3_bn (BatchNormalization) | (None, 56, 56, 256) | 1024 | conv2_block1_3_conv[0][0] |
| conv2_block1_add (Add) | (None, 56, 56, 256) | 0 | conv2_block1_0_bn[0][0] conv2_block1_3_bn[0][0] |
| conv2_block1_out (Activation) | (None, 56, 56, 256) | 0 | conv2_block1_add[0][0] |
| conv2_block2_1_conv (Conv2D) | (None, 56, 56, 64) | 16448 | conv2_block1_out[0][0] |
| conv2_block2_1_bn (BatchNormalization) | (None, 56, 56, 64) | 256 | conv2_block2_1_conv[0][0] |
| conv2_block2_1_relu (Activation) | (None, 56, 56, 64) | 0 | conv2_block2_1_bn[0][0] |
| conv2_block2_2_conv (Conv2D) | (None, 56, 56, 64) | 36928 | conv2_block2_1_relu[0][0] |
| conv2_block2_2_bn (BatchNormalization) | (None, 56, 56, 64) | 256 | conv2_block2_2_conv[0][0] |
| conv2_block2_2_relu (Activation) | (None, 56, 56, 64) | 0 | conv2_block2_2_bn[0][0] |
| conv2_block2_3_conv (Conv2D) | (None, 56, 56, 256) | 16640 | conv2_block2_2_relu[0][0] |
| conv2_block2_3_bn (BatchNormalization) | (None, 56, 56, 256) | 1024 | conv2_block2_3_conv[0][0] |
| conv2_block2_add (Add) | (None, 56, 56, 256) | 0 | conv2_block1_out[0][0] conv2_block2_3_bn[0][0] |
| conv2_block2_out (Activation) | (None, 56, 56, 256) | 0 | conv2_block2_add[0][0] |

```
In [74]: from keras.applications.resnet50 import ResNet50, preprocess_input
from PIL import ImageDraw, ImageFont, Image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
image_generator = ImageDataGenerator(rescale = 1./255, preprocessing_function=preprocess_input,
                                     horizontal_flip = True,
                                     width_shift_range = 0.2,
                                     height_shift_range = 0.2,
                                     shear_range = 0.2,
                                     zoom_range = 0.2)

train_gen = image_generator.flow_from_directory(batch_size = 1,
                                              directory= train_path,
                                              shuffle = True,
                                              target_size = (224, 224),
                                              color_mode="rgb",
                                              class_mode = 'categorical',
                                              subset= 'training',
                                              seed = 0)

validation_gen = image_generator.flow_from_directory(batch_size = 1,
                                                    directory= valid_path,
                                                    shuffle = True,
                                                    color_mode="rgb",
                                                    target_size = (224, 224),
                                                    class_mode = 'categorical',
                                                    seed = 0)
```

```
In [79]: from tensorflow.python.keras.models import Sequential, Model
from keras import optimizers
import tensorflow as tf
from keras import layers
from tensorflow.python.keras.layers import Dense, Flatten, GlobalAveragePooling2D, Input, Dropout, BatchNormalization
from tensorflow.python.keras.preprocessing.image import ImageDataGenerator, img_to_array, load_img
basemodel = ResNet50(include_top=False, input_tensor=Input(shape=(224, 224, 3)))
for layer in basemodel.layers:
    layer.trainable=False
result = GlobalAveragePooling2D()(basemodel.output)
# add a fully-connected layer
result = Dense(512, activation='relu')(result)
predictions = Dense(2, activation='sigmoid')(result)
```

The above code builds the Resnet model and creates training and validation data sets. Sequential class is imported which allows one to easily construct a neural network by just stacking layer one after another. GAP (Global Average Pooling) is used to minimize the overfitting by reducing the total number of parameters in the model.


Now model is compiled and fitted using epoch-20 and callbacks-checkpoint. The model is optimized using adam algorithm which combines the best properties of the AdaGrad and RMSProp algorithms to provide an optimization algorithm that can handle sparse gradients on noisy problems.


```
In [82]: ▶ model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

In [83]: ▶ from keras.callbacks import ModelCheckpoint
mc = ModelCheckpoint('best_model.h5', monitor='val_accuracy', mode='max', verbose=1, save_best_only=True)

In [84]: ▶ model.fit_generator(
    train_gen, steps_per_epoch=len(train_gen)/18, shuffle = True,
    epochs=20, validation_data = validation_gen, validation_steps=len(validation_gen)/18, callbacks=[mc])
```

Plot of Model Accuracy and Model loss:

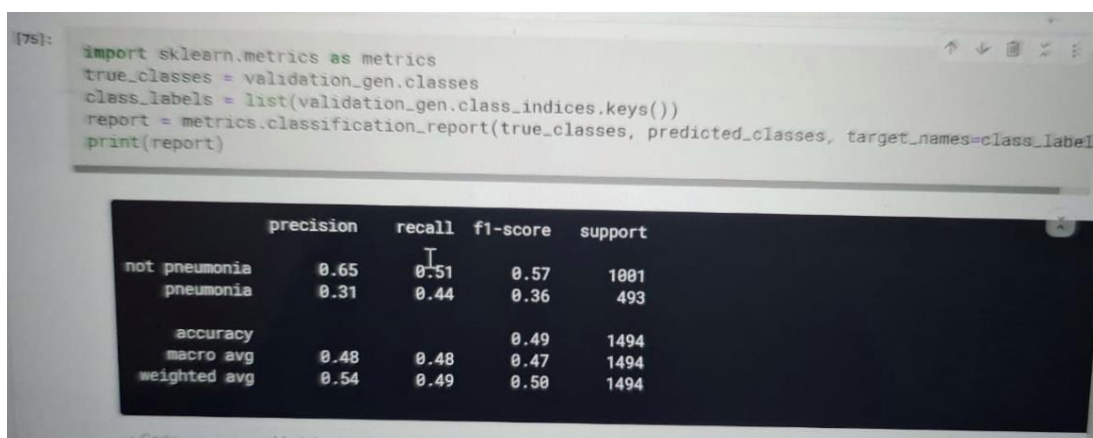




As the epoch is increased, the fluctuation/noise in the validation set gradually reduces and thereafter it is able to predict with moreover consistent accuracy of around 67%. Whereas the maximum possible accuracy observed was 71%.

The training loss is moreover constant which implies model is fully learnt and no further improvement is expected.

The confusion matrix was plotted which revealed a precision 0.65, 0.31 for classes 0 and 1, recall 0.51 and 0.44 for classes 0 and 1 respectively.



We also trained the model with combinations of optimizers and loss (among ‘adam’, ‘sgd’; loss being ‘categorical_crossentropy’ and binary_crossentropy) for few epochs. The accuracy for validation set was compared and the best combination of ‘adam’ optimizer with loss being ‘categorical_crossentropy’ was employed.

```
In [44]: model.compile(loss='binary_crossentropy',
                    optimizer='adam',
                    metrics=['accuracy'])

In [45]: model.fit(
    train_gen, steps_per_epoch=len(train_gen)/18,
    epochs=2, validation_data = validation_gen, validation_steps=len(validation_gen)/2)

Epoch 1/2
1186/1186 [=====] - 559s 472ms/step - loss: 0.6106 - accuracy:
0.7487 - val_loss: 0.5242 - val_accuracy: 0.7721
Epoch 2/2
1186/1186 [=====] - 559s 472ms/step - loss: 0.6149 - accuracy:
0.7015 - val_loss: 0.5345 - val_accuracy: 0.7822

Out[45]: <tensorflow.python.keras.callbacks.History at 0x7fc5a00efe50>
```

```
In [41]: model.compile(loss='categorical_crossentropy',
                    optimizer='sgd',
                    metrics=['accuracy'])

In [42]: model.fit(
    train_gen, steps_per_epoch=len(train_gen)/18,
    epochs=2, validation_data = validation_gen, validation_steps=len(validation_gen)/2)

Epoch 1/2
1186/1186 [=====] - 556s 469ms/step - loss: 3.7190 - accuracy:
0.6627 - val_loss: 3.4335 - val_accuracy: 0.7755
Epoch 2/2
1186/1186 [=====] - 561s 473ms/step - loss: 3.6702 - accuracy:
0.6644 - val_loss: 3.3885 - val_accuracy: 0.7631

Out[42]: <tensorflow.python.keras.callbacks.History at 0x7fc581f01d50>
```

```
In [43]:
```

```

In [38]: model.compile(loss='categorical_crossentropy',
                      optimizer='adam',
                      metrics=['accuracy'])

In [39]: model.fit(
          train_gen, steps_per_epoch=len(train_gen)/18,
          epochs=2, validation_data = validation_gen, validation_steps=len(validation_gen)/2)

Epoch 1/2
1186/1186 [=====] - 555s 468ms/step - loss: 0.5876 - accuracy:
0.7698 - val_loss: 0.6177 - val_accuracy: 0.7729
Epoch 2/2
1186/1186 [=====] - 555s 468ms/step - loss: 0.5636 - accuracy:
0.7648 - val_loss: 0.5701 - val_accuracy: 0.7826

Out[39]: <tensorflow.python.keras.callbacks.History at 0x7fc5e4661550>

In [40]:

```

Summary:

We approached the problem starting with an exploratory data analysis of the CSV files and DICOM images. The training set 9555 belonging to class pneumonia and remaining belonging to normal and non-pneumonic abnormalities.

We tried two models: A mask RCNN model and ResNet based classification model

Mask RCNN model was built with pretrained weights from coco-model weights. On evaluation the bounding box error was around 0.5. A classification model was built with ResNet layers and pretrained weights from 'ImageNet'. The later model was decided to be taken further for finetuning and evaluation.

The model was trained with combinations of optimizers and loss (among 'adam', 'sgd'; loss being 'categorical_crossentropy' and binary_crossentropy)

The parameters were chosen as with 'adam' as optimizer and loss being 'categorical_crossentropy', since it gave the best validation accuracy among the combinations.

The prediction value for validation set was as well calculated and saved to a CSV file.

